

Sentiment Analysis And Other Non-explicit Semantics Of Text

Lu Xiao

lxiao04@syr.edu

213 Hinds Hall



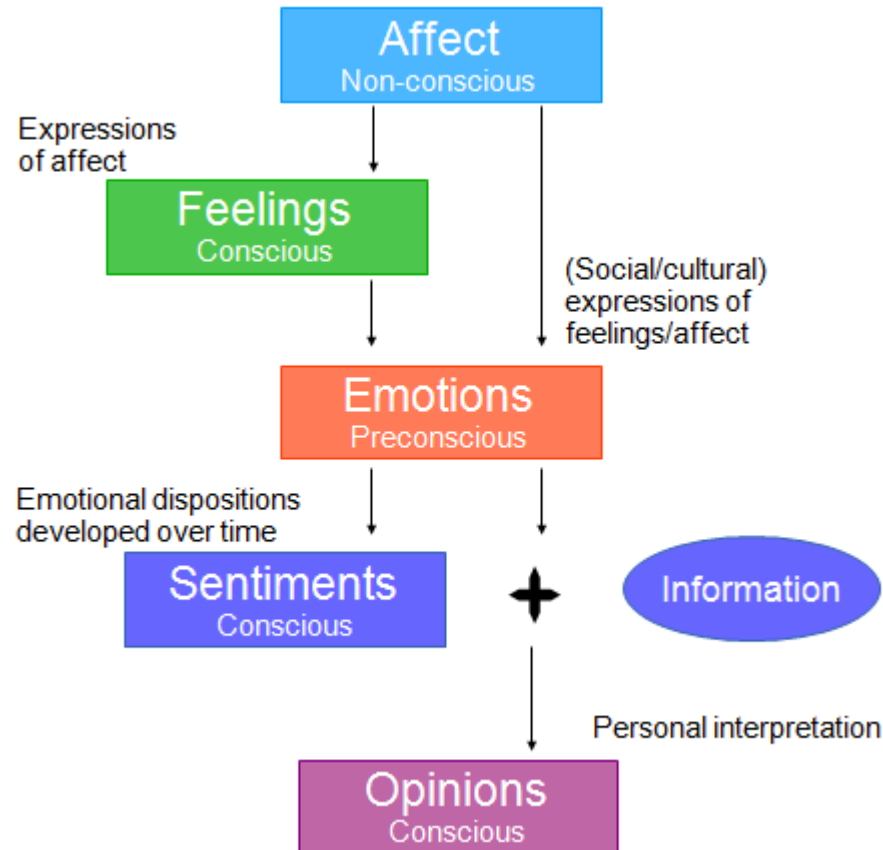
adopted some materials developed in previous courses by Nancy McCracken, Liz Liddy and others; and some instructor resources for the book “Speech and Language Processing” by Daniel Jurafsky and James H. Martin

Affect, Feeling, Emotion, Sentiment, and Opinion Detection (Munezero, Montero, Sutinen, & Pajunen, 2014)

- A major limitation in the automatic detection of affect, feelings, emotions, sentiments, and opinions in text: the lack of proper differentiation between these subjective terms and understanding of how they relate to one another
- The differences between these five subjective terms



Affect, Feeling, Emotion, Sentiment, and Opinion Detection (Munezero, Montero, Sutinen, & Pajunen, 2014)



MUNEZERO, MYRIAM & SUERO MONTERO, CALKIN & SUTINEN, ERKKI & PAJUNEN, JOHN. "ARE THEY DIFFERENT? AFFECT, FEELING, EMOTION, SENTIMENT, AND OPINION DETECTION IN TEXT" AFFECTIVE COMPUTING, IEEE TRANSACTIONS ON. 5. 101-111 (2014)



Scherer Typology Of Affective States

- **Emotion:** brief organically synchronized ... evaluation of a major event
 - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
 - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes:** enduring, affectively colored beliefs, dispositions towards objects or persons
 - *liking, loving, hating, valuing, desiring*
- **Personality traits:** stable personality dispositions and typical behavior tendencies
 - *nervous, anxious, reckless, morose, hostile, jealous*



Sentiment Analysis

- **Emotion:** brief organically synchronized ... evaluation of a major event
 - *angry, sad, joyful, fearful, ashamed, proud, elated*
- **Mood:** diffuse non-caused low-intensity long-duration change in subjective feeling
 - *cheerful, gloomy, irritable, listless, depressed, buoyant*
- **Interpersonal stances:** affective stance toward another person in a specific interaction
 - *friendly, flirtatious, distant, cold, warm, supportive, contemptuous*
- **Attitudes: enduring, affectively colored beliefs, dispositions towards objects or persons**
 - *liking, loving, hating, valuing, desiring*
- **Personality traits:** stable personality dispositions and typical behavior tendencies
 - *nervous, anxious, reckless, morose, hostile, jealous*



Why Sentiment Analysis?

- *Movie*: is this review positive or negative?
- *Products*: what do people think about the new iPhone?
- *Customer relations*: what do people say about your company?
- *Public sentiment*: how is consumer confidence? Is despair increasing?
- *Politics*: what do people think about this candidate or issue?
- *Prediction*: predict election outcomes or market trends from sentiment

Sentiment Analysis

- Sentiment analysis is the detection of attitudes
“enduring, affectively colored beliefs, dispositions
towards objects or persons”
 - **Holder (source)** of attitude
 - **Target (aspect)** of attitude
 - **Type** of attitude
 - From a set of types
 - Like, love, hate, value, desire, etc.
 - Or (more commonly) simple weighted polarity:
 - Positive, negative, neutral, together with strength
 - **Text** containing the attitude
 - Sentence or entire document

Sentiment Analysis Task Levels

- **Simplest task is polarity:**
 - Is the attitude of this text positive or negative?
 - Negative / positive attitude of reporter / blogger
 - Favorable / unfavorable review of a product
 - Right / left political leaning of speaker
- **More complex:**
 - Rank the attitude of this text from 1 to 5
 - Sometimes called strength or intensity
- **Advanced:**
 - Detect the target, source or complex attitude type
 - May also be referred to as
opinion extraction, opinion mining, or sentiment mining



Challenges – Word Sense Ambiguity

- Consider classifying a subjective text unit as either positive or negative.
 - Example: The most thoroughly joyless and inept film of the year, and one of the worst of the decade. [Mick LaSalle, describing *Gigli*]
- Can't we just look for words like *great* or *terrible* ?
 - Yes, but ...
 - Words have different meanings in different contexts

Challenges – Subtlety in Language

- Subtlety, sarcasm or metaphor:
 - Perfume review in *Perfumes: the Guide*:
 - “If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut.”
 - Dorothy Parker on Katherine Hepburn
 - “She runs the gamut of emotions from A to B”



Challenges - Thwarted Expectations And Ordering Effects

- “This film should be brilliant. It sounds like a great plot, the actors are first grade, and the supporting cast is good as well, and Stallone is attempting to deliver a good performance. However, it **can't hold up.**”
- Well as usual Keanu Reeves is nothing special, but surprisingly, the very talented Laurence Fishbourne is **not so good** either, I was surprised.

Challenges - Domain Adaptation

- Certain sentiment-related indicators seem domain-dependent.
 - *.Read the book..*: good for book reviews, bad for movie reviews
 - *.Unpredictable..*: good for movie plots, bad for a car's steering [Turney '02]
- In general, sentiment classifiers (especially those created via supervised learning) have been shown to often be domain dependent
 - [Turney '02, Engström '04, Read 05, Aue & Gamon '05, Blitzer, Dredze & Pereira '07].

Sentiment Detection: Polarity And Intensity



Sentiment Polarity

- Classic Sentiment polarity task from Pang and Lee:
 - Is an IMDB movie review positive or negative?
 - Data: *Polarity Data 2.0*: (people indicate polarity of own review)
 - <http://www.cs.cornell.edu/people/pabo/movie-review-data>
- **Treat as a document classification task**
 - Positive, negative, and (possibly) neutral
 - sentiment words are often more important than topic words, e.g., great, excellent, horrible, bad, worst, etc.

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. EMNLP-2002, 79—86.

Bo Pang and Lillian Lee. 2004. A Sentimental Education: Sentiment Analysis Using Subjectivity Summarization Based on Minimum Cuts. ACL, 271-278

Treat As A Classification Problem

- **Tokenization**
 - May be some differences, especially for text from social media
- **Feature Extraction**
 - The most important part!
- **Classification using different classifiers**
 - Naïve Bayes
 - MaxEnt
 - SVM
 - Deep learning



Sentiment Tokenization Issues

- For text from web, deal with HTML and XML markup
- Or Twitter mark-up (names, hash tags)
- Capitalization (preserve for

Potts emoticons

words in all caps)

- Phone numbers, dates
- Emoticons

[<>]?	# optional hat/brow
[:;=8]	# eyes
[\-o*\']?	# optional nose
[\\)\]\(\[dDpP/\:~\}\{@\ \\\]	# mouth
	#### reverse orientation
[\\)\]\(\[dDpP/\:~\}\{@\ \\\]	# mouth
[\-o*\']?	# optional nose
[:;=8]	# eyes
[<>]?	# optional hat/brow

- Useful code:
 - [Christopher Potts sentiment tokenizer](#)

Extracting Features For Sentiment Classification

- Which words to use?
 - Only adjectives
 - All words
 - All words turns out to work better, at least on this data
- Syntax is not used as often
 - Constituent or dependency parses are occasionally used
 - Particularly at phrase level to find dependencies of opinion words
 - Also for finding the scope of negation
 - Can be used to shift the “valence”
 - For negation, intensification and diminution
 - Very good, deeply suspicious
 - Should have been good
 - He is a great actor, *however* this performance . . .
 - However changes the valence of great to be negative

Handling Negation Is Important!

- How to handle negation:
 - I **didn't** like this movie vs I really like this movie
 - Pang and Lee simple approximation to negation:
 - Add NOT_ to every word between negation and following punctuation:

didn't like this movie , but I



didn't NOT_like NOT_this NOT_movie but I

- Negation has both scope and focus
 - These may be represented in more complex structures
 - Details in Wilson “Fine-grained sentiment analysis”

Sentiment Lexicons

- One of the early approaches to sentiment analysis was to just count the words in each document that had either a positive or negative polarity from a (hand-built) sentiment lexicon.
 - This approach usually not very accurate on individual documents, but it's easy because doesn't need training data.
 - May be useful over aggregate collections or to show trends over time.
- Now we use either presence or frequencies of sentiment words as features of the classifier

MPQA Subjectivity Cues Lexicon

- Gives a list of words that have been judged to be weakly or strongly positive, negative or neutral in subjectivity

- Home page:

http://www.cs.pitt.edu/mpqa/subj_lexicon.html

- 6885 words from 8221 lemmas

- 2718 positive, 4912 negative

- GNU GPL license

- Examples:

type=weaksubj len=1 word1=abandoned pos1=adj stemmed1=n priorpolarity=negative

type=weaksubj len=1 word1=abandonment pos1=noun stemmed1=n priorpolarity=negative

type=weaksubj len=1 word1=abandon pos1=verb stemmed1=y priorpolarity=negative

type=strongsubj len=1 word1=abase pos1=verb stemmed1=y priorpolarity=negative

type=strongsubj len=1 word1=abacement pos1=anypos stemmed1=y priorpolarity=negative

LIWC (Linguistic Inquiry And Word Count)

- Text analysis software based on dictionaries of word dimensions
- Dimensions can be syntactic
 - Pronouns, past-tense verbs
- Dimensions can be semantic
 - Social words, affect, cognitive mechanisms
- Other categories (<http://www.liwc.net/comparedicts.php>)
- James Pennebaker, Univ. of Texas at Austin
 - <http://www.liwc.net/>
- Often used for positive and negative emotion words in opinion mining

ANEW (Affective Norms For English Words)

- From the NIMH Center for the Study of Emotion and Attention at the University of Florida
 - <http://csea.phhp.ufl.edu/Media.html>
 - See also the paper by Dodds and Danforth on Happiness of Large-Scale Written Expressions
 - Free for research use
- Provides a set of emotional ratings for a large number of words in the English language

BING LIU OPINION LEXICON

- [Bing Liu's Page on Opinion Mining](#)
- <http://www.cs.uic.edu/~liub/FBS/opinion-lexicon-English.rar>
- 6786 words
 - 2006 positive
 - 4783 negative

SentiWordNet

- Home page: <http://sentiwordnet.isti.cnr.it/>
- All WordNet synsets automatically annotated for degrees of positivity, negativity, and neutrality/objectiveness
- [estimable(J,3)] “may be computed or estimated”

Pos 0 Neg 0 Obj 1

- [estimable(J,1)] “deserving of respect or high regard”

Pos .75 Neg 0 Obj .25



Which Sentiment Lexicon To Use?

- An area of active research in the sentiment analysis community
- It is now recognized that the amount of overlap between the lexicons is small!
 - But in general, where there is overlap, the sentiment polarity of the words is in agreement, 2% or less disagreement.
 - Except for SentiWordNet, which disagrees up to 25%
 - Chris Potts, Sentiment Symposium Tutorial
- How to represent features from sentiment words:
 - Frequency of all positive and all negative words
 - Presence of positive or negative words (particularly for twitter)
 - Sum of the positive or negative intensity scores
 - ...

Build A Sentiment Lexicon

- For some domains, it has been shown that the best lexicon is one built for that domain
- Automatic lexicon building from unlabeled data
 - Bootstrapping
 - Identify a number of seed words of positive and negative polarity
 - Search for text involving those words that also have connecting words, such as “and”
 - Other words that occur with the connecting word are added to the lexicon with the appropriate polarity

Building A Sentiment Lexicon

- Automatic lexicon building from labeled data
 - In some cases, the domain has lots of text that has been labeled with sentiment
 - Twitter
 - Use tweets labeled with sentiment hashtags: #good, #happy, #bad, #sad
 - Use tweets labeled with happy or sad emoticons
 - Collect words from the positive and negative labeled texts and keep the frequent ones as part of a lexicon
 - Using Mutual Information scores or other measures

VADER

VADER is reported to be among the best sentiment prediction tool available for classifying social media texts and online review comments (Hutto & Gilbert, 2014; Ribeiro, Araújo, Gonçalves, Gonçalves, & Benevenuto, 2016)

- Lexicon
- Syntactic rules, e.g., all capitalization, punctuation marks, etc.

Hutto, C.J., Gilbert, E.E.: VADER: a parsimonious rule-based model for sentiment analysis of social media text. In: Eighth International Conference on Weblogs and Social Media (ICWSM-2014), Ann Arbor, MI, June 2014

Ribeiro, F. N., Araújo, M., Gonçalves, P., Gonçalves, M. A., & Benevenuto, F. (2016). SentiBench-a benchmark comparison of state-of-the-practice sentiment analysis methods. EPJ Data Science, 5(1), 23.



Detection Of Representative Rationales In Wikipedia Afd Deliberations (Mao, Xiao, & Mercer, 2014)

Input a discussion

Vote, deliberation content, and user name

Divide the discussion into groups

Sentence-to-sentence semantic similarity (SEMILAR)

Divide each group into three subgroups (positive, neutral, negative)

Using our recursive algorithm to predict polarity

Extract representative rationales based on the similarity score


Our Recursive Algorithm to Predict Sentiment Polarity at Sentence Level

-Based on the dependency structure tree

- Take a node as input, and the polarity score for the node as output.

-The algorithm assigns a polarity score to each node in the dependency structure tree

-Integrate five types of negations

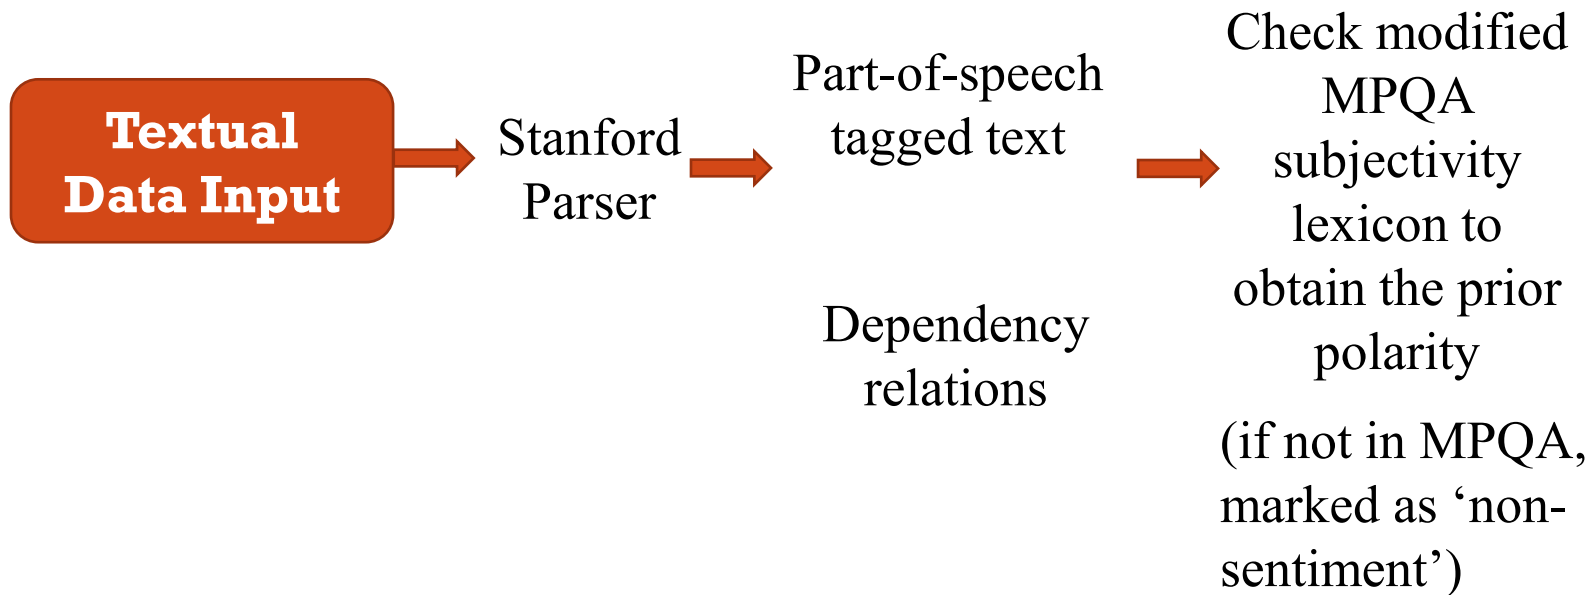


Five Negations

- Local negation: A not usually modifies the sentiment word.
 - “The place is *not notable*.”
- Predicate negation: using verbs with negative polarity.
 - “I *disagree* that the place is *notable*.”
- Subject negation: a subject leads to the negation of its predicate.
 - “*Neither* one of us *agrees* that the place is *notable*.”
- Preposition negation: the polarity of the object following the preposition “of” can be changed by the word modified by the preposition.
 - “It is a *violation* of *notability*.”
- Modifier negation: some sentiment word’s polarity can be negated by its modifier.
 - “The place is of *indeterminable notability*.”


Sentiment Analysis

- Determine the sentiment polarity of a statement in our language context (“notable”)
 - MPQA Subjectivity Lexicon + additional words



MPQA format:

type=strongsubj len=1 word1=aberration pos1=adj stemmed1=n
priorpolarity=negative



Experiment And Evaluation On Sentence Polarity Prediction

- Methods
 - Stanford sentiment analysis tool vs. our algorithm

	Stanford sentiment analysis	recursive algorithm without learning	recursive algorithm with machine learning
Accuracy (%)	48.73	58.47	60.17

- Accuracy of Stanford sentiment analysis tool classifying movie review
 - 5-category: 45.7%, 2-category: 85.4%

Mao, W., Xiao, L., & Mercer, R. (2014). The use of text similarity and sentiment analysis to examine rationales in the large-scale online deliberations. In *Proceedings of the 5th Workshop on Computational Approaches to Subjectivity, Sentiment and Social Media Analysis* (pp. 147-153). , <http://www.aclweb.org/anthology/W14-2624>

Summary On Sentiment Analysis

- Understanding semantics of less factual aspects of text
- Generally modeled as classification or regression task
 - predict a binary label for polarity
 - or predict an ordinal label for the level of sentiment
- Features:
 - Negation is important
 - Using all (filtered) words works well for some tasks
 - Finding subsets of words may help in other tasks
 - Hand-built polarity lexicons
 - Use seeds and semi-supervised learning to induce lexicons

Web Resources In Blackboard



Opinion Analysis



Opinion Mining

- The **third level of sentiment analysis** is sometimes called opinion mining because you are finding sentiment towards aspects or attributes
- Businesses spend a huge amount of money to find consumer sentiments and opinions.
 - Consultants, surveys and focus groups, etc
 - Text in the form of transcripts of interviews or survey responses
- Opinions also available on the web
 - product reviews
 - blogs, discussion groups

Sentence Level Detection

- Sentence level or sub-sentence level detection of subjectivity as a binary classifier
 - Wiebe, many projects
 - Pang and Lee – for movie reviews, first determine which sentences express opinions and then label for opinion polarity
- Clause level opinion strength
 - Wilson, “How mad are you?”
- Detection of sentences with subjectivity or sentiment is important to then find aspects or attributes
 - *The food was great but the service was awful.*

Finding Aspect/Attribute/Target Of Sentiment

- Frequent phrases + rules
 - Find all highly frequent phrases across reviews (“fish tacos”)
 - Filter by rules like “occurs right after sentiment word”
 - “...great fish tacos” means fish tacos a likely aspect

Casino	casino, buffet, pool, resort, beds
Children’s Barber	haircut, job, experience, kids
Greek Restaurant	food, wine, service, appetizer, lamb
Department Store	selection, department, sales, shop, clothing

M. Hu and B. Liu. 2004. Mining and summarizing customer reviews. In Proceedings of KDD.

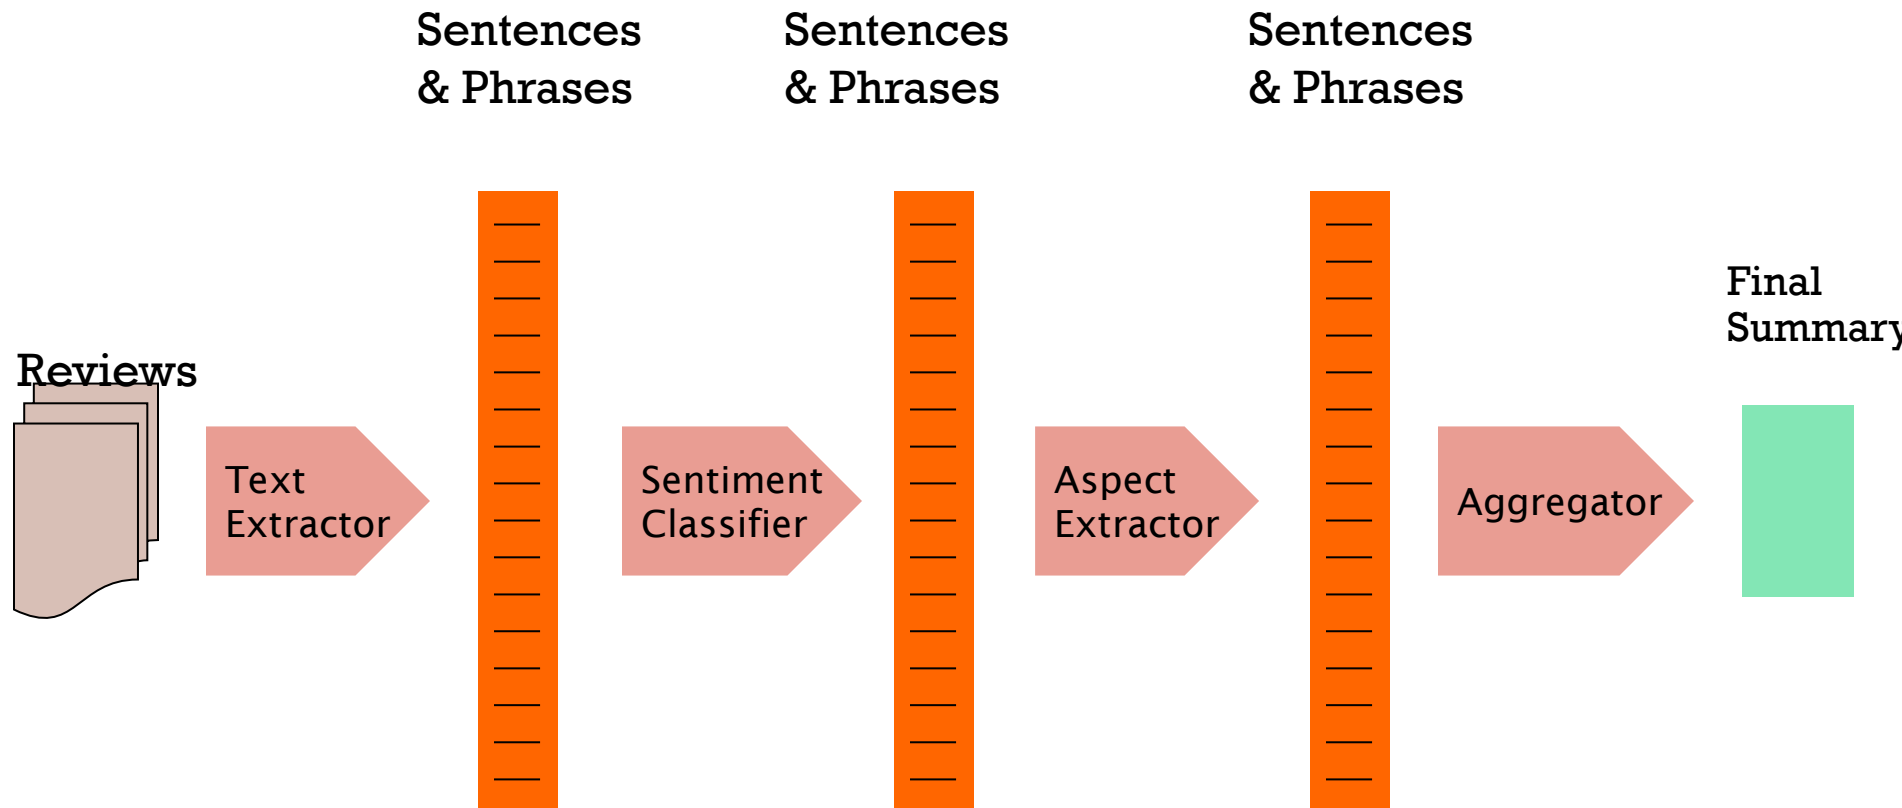
S. Blair-Goldensohn, K. Hannan, R. McDonald, T. Neylon, G. Reis, and J. Reynar. 2008. Building a Sentiment Summarizer for Local Service Reviews. WWW Workshop.



Finding Aspect/Attribute/Target Of Sentiment

- But the aspect name may not be in the sentence
- Other methods to find aspects:
 - For restaurants/hotels, aspects are well-understood
 - Supervised classification
 - Hand-label a small corpus of restaurant review sentences with aspect such as food, décor, service, value, NONE
 - Train a classifier to assign an aspect to a sentence
 - “Given this sentence, is the aspect *food*, *décor*, *service*, *value*, or *NONE*”

Putting It All Together: Finding Sentiment/Opinion For Aspects



Joint Topic/Sentiment Analysis

- An alternative approach to first finding the aspect or attribute and then the opinion or sentiment is to find them both in the same classification
 - Comparative studies of related products
 - Topics that have various features and attributes
 - Consumers
 - Political areas

Example Results For Opinion Of Aspect

Rooms (3/5 stars, 41 comments)

(+) The room was clean and everything worked fine – even the water pressure ...

(+) We went because of the free room and was pleasantly pleased ...

(-) ...the worst hotel I had ever stayed at ...

Service (3/5 stars, 31 comments)

(+) Upon checking out another couple was checking early due to a problem ...

(+) Every single hotel staff member treated us great and answered every ...

(-) The food is cold and the service gives new meaning to SLOW.

Dining (3/5 stars, 18 comments)

(+) our favorite place to stay in biloxi.the food is great also the service ...

(+) Offer of free buffet for joining the Play



Feature-based Summary (Hu AND Liu, Kdd-04)

- From reviews, extract a summary:

GREAT Camera., Jun 3, 2004

Reviewer: **jprice174** from Atlanta, Ga.

I did a lot of research last year before I bought this camera... It kinda hurt to leave behind my beloved nikon 35mm SLR, but I was going to Italy, and I needed something smaller, and digital.

The **pictures** coming out of this camera are amazing. The '**auto**' feature takes great pictures most of the time. And with digital, you're not wasting film if the picture doesn't come out. ...

Feature Based Summary:

Feature1: **pictures**

Positive: 12

- The **pictures** coming out of this camera are amazing.
- Overall this is a good camera with a really good **picture** clarity.

...

Negative: 2

- The **pictures** come out hazy if your hands shake even for a moment during the entire process of taking a picture.
- Focusing on a display rack about 20 feet away in a brightly lit room during day time, **pictures** produced by this camera were blurry and in a shade of orange.

Feature2: **battery life**

...

....